

Atari Breakout with Reinforcement Learning

Table of Contents

- [Atari Breakout with Reinforcement Learning](#)
 - [Table of Contents](#)
 - [Overview](#)
 - [Requirements](#)
 - [Setting Up a Conda Environment](#)
 - [Installing Dependencies](#)
 - [Running the Code](#)
 - [Troubleshooting](#)
 - [Contributing](#)
- [Project badge](#)
 - [Deep Q-learning](#)
- [Resources](#)
 - [Read or watch:](#)
 - [References:](#)
 - [Learning Objectives](#)
- [Requirements](#)
 - [General](#)
- [Tasks](#)
 - [0. Breakout](#)

Overview

This project aims to train a reinforcement learning agent to play Atari's Breakout game. We use Python 3.5 and various libraries like Gym, Keras, and Keras-RL to accomplish this. The project contains two main scripts:

- [train.py](#): Trains the agent using DQN (Deep Q-Network).
- [play.py](#): Allows the trained agent to play the game.

Requirements

- Python 3.5
- NumPy 1.15
- Gym 0.17.2
- Keras 2.2.5
- Keras-RL 0.4.2

Setting Up a Conda Environment

Conda is a package and environment management system that allows you to install software packages and manage different environments for various projects. Follow these steps to set up a Conda environment:

1. **Install Anaconda or Miniconda:** Download from [Anaconda](#) or [Miniconda](#).
2. **Open Terminal:** Open your terminal (Command Prompt on Windows, Terminal on macOS or Linux).
3. **Create a New Environment:** Run `conda create --name atari_breakout python=3.5`.
4. **Activate the Environment:** Run `conda activate atari_breakout` on Windows or `source activate atari_breakout` on macOS and Linux.

Environment Creation/Configuration

Create these files under a folder of your choice.

Environment.yml

```
name: deep
channels:
  - anaconda
  - defaults
dependencies:
  - python=3.6
  - pip
  - pip:
    - pip
```

(notice that I wrote pip three times; the first one is to install pip, but, that installs pip version 20.2.4
the two other lines are to upgrade pip to the latest version to avoid annoying errors/warnings later)

Requirements.txt

```
h5py==2.10.0
keras==2.2.4
keras-rl==0.4.2
numpy==1.18.5
opencv-python==4.4.0.42
pyyaml==5.3.1
six==1.15.0
gym
Pillow
tensorflow==1.14.0
```

Store both these files in a folder and open command prompt (powershell for windows and terminal in ubuntu :)) and cd to the directory where you put those files.

You need to create the anaconda environment, so enter this command:

```
conda env create -f environment.yml
```

Second, you should activate the environment created

```
conda activate deep
```

(notice that the name of the environment deep is mentioned in the first line in the environment.yml file)

So now, that you are inside the anaconda environment. You need to complete the installation of the remaining requirements.

```
pip install -r requirements.txt # installs all packages in the file
```

Lastly, is to install atari_py
Windows:

```
pip install --no-index -f https://github.com/Kojoley/atari-py/releases atari_py
```

For more details, check the [Conda documentation](#).

Installing Dependencies

After activating your Conda environment, install the required packages:

```
conda install numpy=1.15 gym=0.17.2  
pip install keras==2.2.5 keras-rl==0.4.2
```

Running the Code

- Train the Agent: Run `python train.py` to train the agent. The trained model will be saved as `policy.h5`.
- Play the Game: Run `python play.py` to see the trained agent in action.

Troubleshooting

- Conda Command Not Found: Make sure Anaconda/Miniconda is installed and added to your system's PATH. See detailed guide.
- Environment Doesn't Exist: Ensure you have the correct Gym version and have installed the Atari dependencies (`pip install gym[atari]`).

Contributing

Feel free to contribute to this project by opening issues or submitting pull requests.

Project badge

Deep Q-learning

Master

By: Alexa Orrico, Software Engineer at Holberton School

Weight: 6

Manual QA review must be done (request it when you are done with the project)

Resources

Read or watch:

Deep Q-Learning - Combining Neural Networks and Reinforcement Learning
Replay Memory Explained - Experience for Deep Q-Network Training
Training a Deep Q-Network - Reinforcement Learning
Training a Deep Q-Network with Fixed Q-targets - Reinforcement Learning

References:

Setting up anaconda for keras-rl
keras-rl
 rl.policy
 rl.memory
 rl.agents.dqn
Playing Atari with Deep Reinforcement Learning

Learning Objectives

What is Deep Q-learning?
What is the policy network?
What is replay memory?
What is the target network?
Why must we utilize two separate networks during training?
What is keras-rl? How do you use it?

Requirements

General

Allowed editors: vi, vim, emacs

All your files will be interpreted/compiled on Ubuntu 16.04 LTS using python3 (version 3.5)

Your files will be executed with numpy (version 1.15), gym (version 0.17.2), keras (version 2.2.5), and keras-rl (version 0.4.2)

All your files should end with a new line

The first line of all your files should be exactly `#!/usr/bin/env python3`

A README.md file, at the root of the folder of the project, is mandatory

Your code should use the pycodestyle style (version 2.4)

All your modules should have documentation (python3 -c

`'print(__import__("my_module").__doc__)'`)

All your classes should have documentation (python3 -c

`'print(__import__("my_module").MyClass.__doc__)'`)

All your functions (inside and outside a class) should have documentation

(python3 -c `'print(__import__("my_module").my_function.__doc__)'` and python3 -c

`'print(__import__("my_module").MyClass.my_function.__doc__)'`)

All your files must be executable

Your code should use the minimum number of operations

Installing Keras-RL

`pip install --user keras-rl`

Dependencies (that should already be installed)

`pip install --user keras==2.2.4`

`pip install --user Pillow`

`pip install --user h5py`

Tasks

0. Breakout

mandatory

Write a python script `train.py` that utilizes keras, keras-rl, and gym to train an agent that can play Atari's Breakout:

Your script should utilize keras-rl's `DQNAgent`, `SequentialMemory`, and `EpsGreedyQPolicy`

Your script should save the final policy network as `policy.h5`

Write a python script `play.py` that can display a game played by the agent trained by `train.py`:

Your script should load the policy network saved in policy.h5
Your agent should use the GreedyQPolicy

Repo:

GitHub repository: holbertonschool-machine_learning
Directory: reinforcement_learning/deep_q_learning
File: train.py, play.py