

Resources

Read or watch:

An introduction to Reinforcement Learning
Simple Reinforcement Learning: Q-learning
Markov Decision Processes (MDPs) - Structuring a Reinforcement Learning Problem
Expected Return - What Drives a Reinforcement Learning Agent in an MDP
Policies and Value Functions - Good Actions for a Reinforcement Learning Agent
What do Reinforcement Learning Algorithms Learn - Optimal Policies
Q-Learning Explained - A Reinforcement Learning Technique
Exploration vs. Exploitation - Learning the Optimal Reinforcement Learning Policy
OpenAI Gym and Python for Q-learning - Reinforcement Learning Code Project
Train Q-learning Agent with Python - Reinforcement Learning Code Project
Markov Decision Processes

Definitions to skim:

Reinforcement Learning
Markov Decision Process
Q-learning

References:

OpenAI Gym
OpenAI Gym: Frozen Lake env

Learning Objectives

What is a Markov Decision Process?
What is an environment?
What is an agent?
What is a state?
What is a policy function?
What is a value function? a state-value function? an action-value function?
What is a discount factor?
What is the Bellman equation?
What is epsilon greedy?
What is Q-learning?

Requirements

General

Allowed editors: vi, vim, emacs

All your files will be interpreted/compiled on Ubuntu 16.04 LTS using python3 (version 3.5)

Your files will be executed with numpy (version 1.15), and gym (version 0.7)

All your files should end with a new line

The first line of all your files should be exactly `#!/usr/bin/env python3`

A README.md file, at the root of the folder of the project, is mandatory

Your code should use the pycodestyle style (version 2.4)

All your modules should have documentation (python3 -c

`'print(__import__("my_module").__doc__)'`)

All your classes should have documentation (python3 -c

`'print(__import__("my_module").MyClass.__doc__)'`)

All your functions (inside and outside a class) should have documentation

(python3 -c `'print(__import__("my_module").my_function.__doc__)'` and python3 -c

`'print(__import__("my_module").MyClass.my_function.__doc__)'`)

All your files must be executable

Your code should use the minimum number of operations

Installing OpenAI's Gym

```
pip install --user gym
```